

PRODUCT RESEARCH · 7 COMPETITORS · 31 GAPS

Jelly Studio vs. the market

Feature gaps & what to build

DATE

August 16, 2026

PREPARED FOR

Jared

SCOPE

TubeGen · Vidnoz · Argil · StoryShort · NoLang · Zebracat · Faceless.so

WHAT'S INSIDE

- 02 The verdict, Jelly's unique advantages, and the "do not copy" list
- 03 The market on one page — pricing, real cost per minute, free-tier truth, trust signals
- 04 The gap matrix — all 31 gaps, grouped NOW / NEXT / LATER / SKIP, with how we add each one
- 10 Build order for this month — the ten-item checklist
- 11 Seven competitor teardowns — one page each

9

gaps tagged **NOW** — UI over engines that already run, \$0 new spend

9

NEXT, once the first five payers prove the loop

8

LATER — teams, API, analytics, growth loops

5

deliberate **SKIPS** — avatars, stock, autopilot, model menus, subscriptions

The short version. Nobody out there beats Jelly on capability — they beat it on *legibility*. Every rival shows the price before the click, drafts cheap before it charges, and refunds a failed render automatically. Jelly does the hard part already, then surprises the user with the bill. This month is cheap UI over engines that are **already running**, not new engines.

The verdict

Jelly already has the hard parts (generated visuals, cloned voices, locked style, YouTube publish, receipts) and most rivals win on legibility and convenience rather than capability, so the month should go to cheap UI over existing engines: show the price before the click, make stills-draft the default with motion as an opt-in second spend, refund failures automatically, add chapters, per-scene motion toggles, overlay toggles, Remix, a pricing calculator, BYO narration and scheduled publish. Skip avatars, stock footage, autopilot series, model menus and subscriptions on purpose and say so on the landing page, because those are the sources of every competitor's 2-star reviews and Jelly's absence of them is the pitch. Multi-language, 9:16 cuts, thumbnail A/B, longer exports and an API are real gaps but belong after the first five payers prove the core loop.

What Jelly has that none of them do

- True pay-per-render at ~\$1-7 per long-form video with an itemized receipt, versus \$39-\$849/mo subscriptions with expiring credits everywhere else.
- No stock footage, no avatars: every frame is generated in a locked art style with consistent characters, which is exactly what rivals' users complain they cannot get.
- IndexTTS-2 cloned voices with a Voice Tuner and consent-gated uploads, cheaper than ElevenLabs-metered narration on every competitor.
- Kimi planner + Gemini scene QA + a 137-rule production doctrine means output is reviewed, not sprayed, which is the monetization-safe posture post July 2025.
- Own GPU pipeline (DGX + Modal) so cost per minute is controlled and can be shown honestly before render.
- Long-form 16:9 first with a 7-step editor and YouTube OAuth publish; most rivals are 30-90 s vertical tools bolted onto long-form claims.
- Invite-only beta with view-as, feedback and changelog already wired, so the 5 target users can be hand-held into paying.

Do **NOT** copy this

- ✗ Autopilot series that generate and post daily unattended — it is the exact 'AI slop' pattern YouTube's inauthentic-content policy demonetizes and it removes the human review Jelly's QA depends on.
- ✗ Stock B-roll as default fill (Pexels/Storyblocks) — reviewers of TubeGen, StoryShort, Zebracat and NoLang all cite the generic stock look; Jelly's generated-only visuals are the differentiator.
- ✗ Talking avatars / lip-synced presenters — every rival gets 'uncanny lip-sync' complaints and faceless creators do not want a face.
- ✗ Monthly subscriptions with expiring credits, no refunds and 'billed yearly' traps — Vidnoz (2.3/5), StoryShort (2.5/5), Zebracat and Argil reviews are dominated by billing anger; pay-per-render is Jelly's trust edge.
- ✗ Frontier-model menus (Veo/Kling/Sora per clip) — breaks locked style, explodes per-video cost past \$7, and shifts the product from channel tool to model playground.
- ✗ Gray-hat growth loops (warmed accounts, crypto tokens, spin-and-win, review-for-credits) — cheap trust signals that read as spam.
- ✗ Live Niche Finder / channel scraping — TubeGen's is called thin by every reviewer and it is a separate product from rendering.
- ✗ Dual metering (video caps plus AI credits) or per-video 'free edit' counts — Zebracat's biggest complaint; keep one number, dollars.

THE THREE LINES THAT BELONG ON THE LANDING PAGE

01

No subscription. You pay per render — about **\$1-7** for a finished long-form video, itemized on a receipt. Nothing expires.

02

No stock footage. No avatars. Every frame is generated in your locked art style, with the same characters across the whole channel.

03

See the price before you click Generate — and a render that fails refunds itself, with the reason on the receipt.

Lines 1 and 2 are true today. **Line 3 is not yet** — it is items **1** and **3** of the build order on page 10. Close those two and the landing page can make all three claims without an asterisk.

The market, on one page

Every competitor charges a subscription and meters credits that expire. Six of the seven have a public trust problem, and in five of those the complaints are about **billing**, not output quality. That is the hole Jelly is already standing in.

COMPETITOR	PRICING MODEL	EFFECTIVE COST	MAX LENGTH	FREE TIER – THE TRUTH	TRUST SIGNAL & HEADLINE COMPLAINT
TubeGen AI tubegen.ai	Subscription + credits. Starter \$149 / Pro \$297 / Premium \$849 per mo. À-la-carte “Build Your Own Plan” suites from \$27/mo. No refunds.	\$1.94–2.47/min their own “typical” figure	30 min export cap	None. No free trial at all — you buy blind at \$149.	Trustpilot 4.1/5 from just 18 reviews — 12 posted on one day (23 Apr 2026). “Best” images burn 90 or even when the batch comes out wrong; refunds are manual credits via Discord.
Vidnoz AI vidnoz.com	Freemium + subscription + add-ons, credit-metered. Starter \$26.99/mo (15 min/mo), Business \$74.99/mo (30 min/mo). Voice Clone +\$9.99/mo, Avatar Pro +\$299/yr.	~\$1.80–2.50/min of plan minutes, add-ons on top	3 min free · 60 min paid	Real: 60 credits/day, no card — but 720p, watermarked, translator capped at 60–90 s.	Trustpilot ~2.3/5, ~42% one-star — almost entirely billing: “billed yearly” fine print, surprise renewals, \$220 charged when expecting \$19.99 , no refunds on monthly.
Argil argil.ai	Subscription + credits, no pure pay-per-render. Classic \$39 / Pro \$149 / Scale \$499 per mo. PAYG top-ups never expire; Auto Top-Up has a spend cap.	~\$1.50–5.00/min Atom 160 cr/min = \$4.20–4.70	Short-form; not viable past ~10 min	3 trial generations, watermarked. 5-day trial that reviewers say paywalls at signup.	Trustpilot 2.8/5 (3 reviews) — trial paywalled immediately, no pro-rata refund on annual. Uncanny lip-sync; billing geo-locked (“not available in your region”).
StoryShort AI storyshort.ai	Subscription only, credit-metered. Starter \$39 / Growth \$69 / Influencer \$129 / Ultra \$199 per mo. “All sales are final.”	~\$2–5+/video long-form 85–150 cr, gated to \$129+/mo	Up to 10 min (marketing claims 30)	~10 credits, preview only — export is watermarked/locked. FAQ claims “no watermark”; the app disagrees.	Trustpilot 2.5/5 — emails unanswered, no complaint path, no refunds. An affiliate says >\$1,000 of commissions were marked paid then reversed .
NoLang no-lang.com	Subscription + credits, JPY only . Standard ¥2,980 (~\$20) / Premium ¥7,980 (~\$53) / Business ¥100k+. API is a separate prepaid wallet.	~\$0.70–2.30/min ≈¥150/min on Standard	3 min on Standard; longer = Premium	Real but thin: 3 generations/mo, watermark, videos deleted after 2 weeks .	No Trustpilot / G2 / Product Hunt footprint — zero English reviews. Generic AI-picked stock visuals, TTS mispronunciation, PC-only, BGM triggers YouTube claims.
Zebracat zebracat.ai	Subscription + dual meter (video count AND AI credits). Cat \$39 / Super Cat \$98 / Unlimited \$198 per mo; Enterprise \$398–599+.	Per-component credits AI visuals 5/min, cloned voice 4/min, zoom 8/min	Hard 5 min on every tier · 1080p cap	5 total videos, 30 s, 720p, watermarked, standard voices — no card.	AppSumo 3.9/5 with 18 one-taco reviews — lifetime credits swapped for a 30-video/mo cap. Credits expire monthly, no rollover, refunds “a nightmare”; “could not export 50% of the videos.”
Faceless.so faceless.so	Subscription + credits. Starter \$24/mo (500 cr ≈ 25 videos) → Ultra \$166/mo. Top-ups \$10 per 100 cr. No refunds once one credit is used.	~\$0.96–\$10/video storyboard 20 cr / lite 50 / pro 100	Short-form only — 30/60/90 s, 9:16	“Sign up free”, but their own compare pages say “No free tier” and the demo disables render + upload.	Almost no independent reviews (Groupify 4.75/5 from 4). Sibling brands sit at 1.8–3.9 TrustScore over refunds. ToS bans automated access while selling an API; crypto token, spin-and-win and \$349/mo “warned accounts” read gray-hat.
Jelly Studio tolley.io/animate	Pay-per-render. No subscription, no credits that expire. Compute at cost + \$0.35 per finished minute, itemized on the receipt.	~\$1–7 per long-form video 2–4.5× under TubeGen	9:00 beta cap (raiseable for paid accounts)	Invite-only beta — no public free tier yet. A no-signup demo is a NEXT-priority gap.	No review footprint yet. The counter-positioning is exactly what every row above is getting flamed for: billing .

Read the last column, not the first. Vidnoz sits at 2.3/5 and StoryShort at 2.5/5 with *working products*. The anger is auto-renewals, expiring credits and “all sales are final.” Jelly’s pay-per-render receipt is not a pricing detail — it is the entire wedge, and the landing page should say so in the first screen.

The gap matrix

All 31 gaps, grouped by when to build them. **NOW rows are in pink.** "Jelly today" is what actually exists in the codebase — most NOW items mean surfacing an engine that already runs, not writing one. Effort is **S/M/L**; cost is new recurring spend.

NOW		Ship this month. Cheap UI over engines Jelly already runs — every one closes a complaint a rival is actively getting.		9 gaps
FEATURE · WHO HAS IT	WHY IT MATTERS · WHERE JELLY IS	HOW WE ADD IT	EFFORT · COST	
Pre-render cost estimate shown at the point of choice (per-image tier, motion on/off, per-minute table) TubeGenStoryShortZebracatFaceless.soArgilNoLang	Every competitor makes cost legible before the button; Jelly's receipt arrives after the spend, which is the wrong moment for a first-time paying user. JELLY TODAY partial — check-budget.ts already computes estimateCents for the 402 path and BillingBlock.tsx can display it, but no visible line item per step before Generate	Surface the existing estimateCents from lib/vater/billing/check-budget.ts as an always-on 'Estimated: \$X.XX' row in components/animate/screens/editor/BillingBlock.tsx and VisualsStep.tsx (stills vs +Wan motion delta), plus a static per-minute table on components/animate/screens/browse/PricingScreen.tsx. lib/vater/billing/check-budget.ts, lib/vater/video-cost.ts, components/animate/screens/editor/BillingBlock.tsx, VisualsStep.tsx, browse/PricingScreen.tsx	S \$0 (uses existing cost math)	
Free animatic preview before paid motion (stills + TTS draft, then opt-in Wan pass) ArgilNoLang	Argil and NoLang render a cheap static draft first and only spend expensive credits on the final; it kills the 'I paid and it came out wrong' complaint that dominates TubeGen/Zebracat reviews. JELLY TODAY partial — post_compose_video and post_animate_all_scenes are separate calls in vater.py, but the UI does not frame stills-first as the default flow	In VisualsStep.tsx make 'Render stills draft' the default button, then expose post_animate_all_scenes (vater.py:9679) as a second explicit 'Add motion (\$X)' step with the estimate inline; keep Gemini scene QA gate before motion. ~/content-autopilot/vater.py post_compose_video/post_animate_all_scenes, components/animate/screens/editor/VisualsStep.tsx	S \$0; actually lowers per-video cost	
Bring-your-own voiceover audio (upload narration, scenes/captions/chapters key off it) TubeGenArgilZebracatNoLangFaceless.so	Many faceless creators already record or buy narration; without an upload path Jelly cannot ingest their existing workflow and forces a cloned voice. JELLY TODAY no — VoiceoverStep.tsx only offers catalog/cloned voices; app/api/vater/upload exists but only for voice-clone samples	Add an 'Upload narration (WAV/MP3)' card in components/animate/screens/editor/VoiceoverStep.tsx that posts to app/api/vater/upload/route.ts, then a vater.py branch that skips _run_index tts_modal and runs the existing forced-alignment caption pass on the uploaded file before _schedule_scenes. components/animate/screens/editor/VoiceoverStep.tsx, app/api/vater/upload/route.ts, ~/content-autopilot/vater.py (audio phase + captions alignment)	M \$0 (Whisper/alignment already local)	
Chapter timestamps + hashtags in the generated description TubeGen	YouTube renders 0:00 chapters as clickable, which measurably lifts retention on 8–9 min videos; Jelly already has scene timings so this is nearly free. JELLY TODAY partial — DescriptionStep.tsx generates a description but no chapters/hashtags	In DescriptionStep.tsx read scenes_meta start times from _schedule_scenes output and have the Kimi prompt emit '0:00 Intro' chapter lines plus 3–5 hashtags; pass through lib/vater/youtube-upload.ts unchanged. components/animate/screens/editor/DescriptionStep.tsx, lib/vater/youtube-upload.ts, vater.py _schedule_scenes (5108)	S \$0	

NOW (cont.)				Ship this month. Cheap UI over engines Jelly already runs — every one closes a complaint a rival is actively getting.		9 gaps
FEATURE · WHO HAS IT	WHY IT MATTERS · WHERE JELLY IS	HOW WE ADD IT			EFFORT · COST	
Camera motion and transition controls per scene (pan/zoom alternate, crossfade length, off) TubeGen StoryShort NoLang Zebracat	Cheap Ken-Burns motion is what makes still-based long-form watchable without paying for Wan clips; users expect to toggle it per scene. JELLY TODAY partial — remotion/src/animations/TransitionAnimations.tsx and animation planner exist, but no user-facing per-scene motion/transition picker	Expose a small per-scene control row in VisualsStep.tsx (pan L/R, zoom in/out, still, crossfade 0–2s) that writes into the Remotion props consumed by remotion/src/animations/TransitionAnimations.tsx. components/animate/screens/editor/VisualsStep.tsx, ~/content-autopilot/remotion/src/animations/*, vater.py _plan_scene_animations (3353)			S \$0	
Auto overlays (charts, maps, section headers) surfaced as a user toggle TubeGen	Number-heavy finance/history scripts look richer with charts and maps; TubeGen markets it as a differentiator and Jelly already has the engine. JELLY TODAY partial — vater.py _classify_scenes_for_overlays + remotion ChartScene/MapScene/HeaderScene exist, but not clearly toggled or previewed in the editor	Add three toggles (charts/maps/headers) in VisualsStep.tsx wired to enable_charts/enable_maps/enable_headers on _annotate_scenes_with_overlays, and show an overlay badge on the scene card. vater.py _classify_scenes_for_overlays (6105), remotion/src/scenes/*, components/animate/screens/editor/VisualsStep.tsx			S \$0 (Qwen local)	
Remix / 'same structure, new topic' from an existing project (reference-video loop, saved presets) NoLang Argil TubeGen	Series consistency is the whole game for a channel; one-click 'new video like this one' cuts setup to zero and drives repeat renders. JELLY TODAY partial — locked-style.ts, style-snapshot.ts and youtube/new-from-style route exist, but no 'duplicate this project with new topic' action	Add 'Remix' on browse/ProjectDetail.tsx that calls app/api/vater/youtube/new-from-style with the source project's style snapshot, voice id, aspect, soundtrack and character set pre-filled, dropping the user into ScriptStep with the topic box focused. components/animate/screens/browse/ProjectDetail.tsx, app/api/vater/youtube/new-from-style, lib/vater/style-snapshot.ts			S \$0	
Failed-render auto-refund with plain reason Argil NoLang StoryShort Faceless.so	Trust for a receipts-based product; rivals refund credits automatically within minutes and say why. JELLY TODAY partial — billing/reconcile-renders.ts and ledgers.ts exist; refund on failure is manual	In lib/vater/billing/reconcile-renders.ts credit back the ledger entry when pipeline-status reports failed, and show 'Refunded \$X — reason' in the receipt and System log. lib/vater/billing/reconcile-renders.ts, lib/vater/billing/ledger.ts, app/api/vater/pipeline-status/route.ts			S \$0	
Interactive pricing calculator (videos × length × motion → cost) TubeGen Zebracat	Converts a 'no subscription' pitch into a number a creator can compare to \$149/mo. JELLY TODAY no	Add a 3-slider calculator to components/animate/screens/browse/PricingScreen.tsx and the landing page using DEFAULT_OPS_RATE_PER_MIN and average compute from lib/vater/video-cost.ts. components/animate/screens/browse/PricingScreen.tsx, components/animate/landing/AnimateLanding.tsx, lib/vater/video-cost.ts			S \$0	

NEXT

After the first five payers prove the core loop. Real gaps, but none of them block a first paid render.

9 gaps

FEATURE · WHO HAS IT	WHY IT MATTERS · WHERE JELLY IS	HOW WE ADD IT	EFFORT · COST
Multi-language narration (same script re-voiced in other languages) TubeGen Vidnoz Argil StoryShort NoLang Zebracat Faceless.so	Regional-channel duplication is a top faceless-creator play and every rival ships 8–175 languages; Jelly is English-only in practice. JELLY TODAY partial — vater.py carries a 'language' field (default 'en') and an ElevenLabs voices route already exists (app/api/vater/voices/elevenlabs) , but no language picker and IndexTTS-2 clones are English-tuned	Add a language select in VoiceoverStep.tsx that sets meta.language, route non-English jobs to the existing ElevenLabs path (multilingual v2) via app/api/vater/voices/elevenlabs, and add a Kimi 'translate script' action in ScriptStep.tsx. components/animate/screens/editor/VoiceoverStep.tsx, ScriptStep.tsx, app/api/vater/voices/elevenlabs/route.ts, vater.py language plumbing (line ~3600)	M ElevenLabs ~\$0.15–0.30/min metered per video
9:16 vertical output and post-hoc aspect conversion of an existing project Vidnoz Argil StoryShort NoLang Zebracat Faceless.so	Long-form creators repurpose to Shorts for discovery; a one-click 9:16 cut of a finished project is cheap retention. JELLY TODAY partial — _normalize_aspect/_aspect_image_dims exist in vater.py and an aspect setting exists, but no re-render-to-9:16 flow and Remotion compositions default 16:9	Add 'Make a 9:16 cut' on browse/ProjectDetail.tsx that re-runs _generate_scenes_worker with aspect=9:16 reusing narration and captions, and register a vertical composition in ~/content-autopilot/remotion/src/compositions. ~/content-autopilot/vater.py _normalize_aspect (4383), remotion/src/compositions, components/animate/screens/browse/ProjectDetail.tsx	M ~\$0.50–1.50 per re-render (image regen only)
URL / article / PDF / transcript intake to script Vidnoz Argil StoryShort NoLang Zebracat Faceless.so	Creators repurpose blogs, docs and old videos; paste-a-URL is the lowest-friction way to get to a script. JELLY TODAY partial — ScriptStep.tsx accepts topic/transcript, app/api/vater/topic and rss routes exist, no URL/PDF fetch	Add a 'From URL/PDF' input in ScriptStep.tsx that calls a new app/api/vater/topic?url= branch using the existing lib/vater/rss-parser.ts fetch plus pdf-parse, then feeds text into the Kimi script generator. components/animate/screens/editor/ScriptStep.tsx, app/api/vater/topic/route.ts, lib/vater/rss-parser.ts	S \$0
No-signup interactive demo / sample gallery on the landing page TubeGen Faceless.so NoLang Vidnoz	Jelly is invite-only with no trial; a canned walkthrough is the only way a stranger can feel the product before asking for a code. JELLY TODAY partial — lib/vater/demo-data.ts and demo-videos.ts exist and view-as mode exists, but landing has no clickable demo	Add a read-only 'Try the editor' route reusing ProjectShell.tsx with demo-data.ts and render disabled, linked from components/animate/landing/AnimateLanding.tsx next to the invite form. components/animate/landing/AnimateLanding.tsx, lib/vater/demo-data.ts, lib/vater/demo-videos.ts, components/animate/screens/editor/ProjectShell.tsx	M \$0
Thumbnail A/B variants with side-by-side compare (optionally scored) TubeGen	CTR is the top lever for faceless channels and creators already generate 3 thumbnails elsewhere; keeping them in Jelly keeps the publish step in Jelly. JELLY TODAY partial — ThumbnailStep.tsx generates one thumbnail via post_thumbnail; no variants/compare	Let ThumbnailStep.tsx call post_thumbnail (vater.py:10321) 2–3 times with text/expression variations, show a compare grid, and pass the chosen one to lib/vater/youtube-upload.ts; a Gemini 'which is clearer' score is optional. components/animate/screens/editor/ThumbnailStep.tsx, vater.py post_thumbnail, lib/vater/youtube-upload.ts	S ~\$0.05–0.15 extra per video

NEXT (cont.)

After the first five payers prove the core loop. Real gaps, but none of them block a first paid render.

9 gaps

FEATURE · WHO HAS IT	WHY IT MATTERS · WHERE JELLY IS	HOW WE ADD IT	EFFORT · COST
Longer exports (up to 20–30 min) beyond the 9:00 beta cap TubeGenStoryShortVidnoz	Mid-roll ad eligibility needs 8+ min and many documentary channels run 15–25 min; TubeGen exports 30 min, StoryShort claims 30. JELLY TODAY partial — 9:00 cap enforced in lib/vater/script-limits.ts / check-trial-caps.ts	Raise the cap for paid-credit accounts in lib/vater/script-limits.ts and lib/vater/billing/check-trial-caps.ts once the Remotion render (vater.py _render_remotion) is validated at 20 min on Modal, keeping the estimate visible. lib/vater/script-limits.ts, lib/vater/billing/check-trial-caps.ts, vater.py _render_remotion (5198)	M linear render cost, no new spend
Scheduling / calendar for YouTube publish (publishAt) and multi-platform posting StoryShortFaceless.soNoLangArgil	Creators batch-render then drip-post; YouTube’s API supports a scheduled publishAt on the same OAuth Jelly already has, which is a one-field win. JELLY TODAY partial — YouTube OAuth publish is live; social-accounts route has stubs for other platforms; no scheduler	Add a 'Publish at' datetime to components/animate/screens/review/PublishPanel.tsx passing status.publishAt with privacyStatus=private in lib/vater/youtube-upload.ts; defer TikTok/IG (Postiz/Blotato) until users ask. components/animate/screens/review/PublishPanel.tsx, lib/vater/youtube-upload.ts, app/api/vater/youtube/route.ts	S \$0
Caption style presets (fonts, highlight color, position, emoji, karaoke word emphasis) StoryShortFaceless.soArgilZebracatTubeGen	Captions are the most visible styling element and Jelly already has per-word alignment, so presets are just UI over existing data. JELLY TODAY partial — per-word karaoke captions exist; no user-facing style picker	Add 4–6 caption presets (font/size/position/highlight) selectable in VisualsStep.tsx or a Captions sub-card, saved on the project and consumed by remotion TextAnimations. components/animate/screens/editor/VisualsStep.tsx, remotion/src/animations/TextAnimations.tsx, vater.py caption props (5112)	S \$0
Pronunciation dictionary / per-word fix and re-voice one line NoLangVidnozArgil	Cloned voices mispronounce names and numbers; users otherwise re-render the whole narration. JELLY TODAY partial — vater.py has a spoken-vs-written text pass (line ~4100) but no user dictionary or single-line regen	Add a 'Say it as...' pronunciation map on the project in VoiceoverStep.tsx feeding the existing spoken-text normalizer, plus a per-scene 're-voice this line' calling the IndexTTS chunk for that scene only. components/animate/screens/editor/VoiceoverStep.tsx, vater.py spoken-text normalizer (~4100), _run_indextts_modal (3773)	M pennies per re-voice

LATER

Only matters at scale — agencies, teams, growth loops. Building these before revenue is building the wrong company.

8 gaps

FEATURE · WHO HAS IT	WHY IT MATTERS · WHERE JELLY IS	HOW WE ADD IT	EFFORT · COST
Cross-platform analytics dashboard (views/likes per video) StoryShortFaceless.soVidnoz	Nice to have, but creators live in YouTube Studio; a per-video 'views since publish' line next to cost is enough to show ROI. JELLY TODAY partial — AnalyticsScreen.tsx shows cost subview; lib/vater/youtube-status.ts polls upload state	Extend lib/vater/youtube-status.ts to pull videos.list statistics for published ids and show views next to the receipt in AnalyticsScreen.tsx and ProjectDetail.tsx. lib/vater/youtube-status.ts, components/animate/screens/live/AnalyticsScreen.tsx, browse/ProjectDetail.tsx	S \$0 (YouTube Data API quota)

LATER (cont.)		Only matters at scale — agencies, teams, growth loops. Building these before revenue is building the wrong company.	8 gaps
FEATURE · WHO HAS IT	WHY IT MATTERS · WHERE JELLY IS	HOW WE ADD IT	EFFORT · COST
Public REST API + webhooks + MCP/CLI TubeGenArgilStoryShortNoLangFaceless.soVidnoz	Agencies and agent-driven users want to script renders, and Faceless.so shows MCP is a cheap differentiator; but nobody in the first 5 payers will need it. JELLY TODAY no — autopilot-client.ts and app/api/vater/** are session/proxy gated	Later: issue scoped API keys via lib/vater/direct-auth.ts, expose POST /api/vater/youtube/from-script and pipeline-status with bearer auth, and publish an llms.txt; MCP wrapper over the same routes. lib/vater/direct-auth.ts, app/api/vater/youtube/from-script, app/api/vater/pipeline-status	L \$0
Team seats / shared workspaces TubeGenVidnozArgilFaceless.soZebracatNoLang	Only matters once an agency buys; solo faceless creators do not need it. JELLY TODAY no — project-access.ts and job-ownership.ts are single-owner	Later: extend lib/vater/project-access.ts with an org id and invite via beta-invites.ts. lib/vater/project-access.ts, lib/vater/job-ownership.ts, lib/vater/beta-invites.ts	L \$0
Brand kit (logo, colors, fonts auto-applied to captions/thumbnails/end card) VidnozZebracatFaceless.so	Channel recognizability across videos; cheap once you have a locked style object. JELLY TODAY partial — locked style covers art style; no logo/caption font/colors	Add logo + caption font/color fields to the style object in lib/vater/locked-style.ts and StyleWizardModal.tsx and read them in Remotion caption/thumbnail props. lib/vater/locked-style.ts, components/animate/screens/dashboard/StyleWizardModal.tsx, remotion/src/animations/TextAnimations.tsx	S \$0
Per-segment generated music moods (score changes with narration) TubeGenArgil	A single loop for 9 minutes reads amateur; mood segments are the cheapest 'cinematic' upgrade. JELLY TODAY partial — soundtrack/SFX catalog and per-scene SFX injection exist; one track per video	In SoundtrackStep.tsx let the Kimi planner split the script into 2–4 mood segments and pick a catalog track per segment with crossfades in Remotion, no generative music needed. components/animate/screens/editor/SoundtrackStep.tsx, app/api/vater/music-catalog/route.ts, vater.py get_music_catalog (10026)	S \$0
Prompt-a-voice / voice design and larger stock voice library ArgilVidnozStoryShortFaceless.so	Users without a consented sample want a distinctive narrator fast; a bigger curated library plus optional ElevenLabs picker covers it. JELLY TODAY partial — shared IndexTTS library + ElevenLabs popular-voices route exist; no voice-design	Expose the existing app/api/vater/voices/elevenlabs list in Voices.tsx as an 'ElevenLabs (metered)' tab and add 6–10 more curated IndexTTS presets; skip voice-design. components/animate/screens/studio/Voices.tsx, app/api/vater/voices/elevenlabs/route.ts, lib/vater/voice-catalog.ts	S ElevenLabs metered per minute
Character library page (reuse named characters across projects, re-render in any style) TubeGenStoryShort	Recurring hosts across a channel are Jelly's strength but characters are buried inside the style flow. JELLY TODAY partial — consistent characters + _generate_character_descriptor exist; no standalone browse/manage screen	Add a Characters list under components/animate/screens/browse (mirror StylesListEmbed.tsx) reading character descriptors and letting users pick them in StyleWizardModal.tsx. components/animate/screens/browse/StylesListEmbed.tsx (pattern), dashboard/StyleWizardModal.tsx, vater.py _generate_character_descriptor (10568)	S \$0

LATER (cont.)				Only matters at scale — agencies, teams, growth loops. Building these before revenue is building the wrong company.	8 gaps
FEATURE · WHO HAS IT	WHY IT MATTERS · WHERE JELLY IS	HOW WE ADD IT	EFFORT · COST		
Community + roadmap voting + affiliate/referral TubeGen Zebracat Faceless.so StoryShort Argil	Growth loops matter after product-market fit; for 5 payers a personal Telegram/Discord thread beats a program. JELLY TODAY partial — feedback + changelog exist; no referral	Later: give each user 2 invite codes via lib/vater/beta-invites.ts with a \$5 credit on the invitee's first render (lib/vater/credit-packs.ts). lib/vater/beta-invites.ts, lib/vater/credit-packs.ts	S \$5 credit per referral		
SKIP — ON PURPOSE				Listed so nobody re-opens the argument. These are the sources of every competitor's 2-star reviews. Say "no" on the landing page.	5 gaps
FEATURE · WHO HAS IT	WHY IT MATTERS · WHERE JELLY IS	HOW WE ADD IT	EFFORT · COST		
Stock B-roll mix (Pexels/Pixabay/Storyblocks) with ratio slider TubeGen StoryShort Zebracat Faceless.so Argil NoLang	Real footage helps news/true-crime niches, but it is also the source of the 'generic AI slop' complaint every rival gets. JELLY TODAY no — deliberate 'no stock footage' stance	If ever added, wire a Pexels/Pixabay search in VisualsStep.tsx as a per-scene 'replace with footage' option only, never as the default fill. components/animate/screens/editor/VisualsStep.tsx, Pexels API (free), Pixabay API (free)	M \$0 API, but brand cost		
Talking avatars / lip-synced presenters / digital twins Vidnoz Argil Zebracat StoryShort NoLang TubeGen Faceless.so	Avatars are the default pitch of half the market, but faceless long-form creators explicitly do not want a face and users report uncanny lip-sync everywhere. JELLY TODAY no — Animate2 driver clips exist for characters, not talking heads	Do not build; at most keep Animate2 character-motion clips as the 'presenter' analogue. n/a	L HeyGen-style API \$\$\$ per minute		
Autopilot series that auto-generates and posts daily StoryShort Faceless.so	'Channel on autopilot' sells, but it produces the spam that YouTube's inauthentic-content policy targets and it removes the human review that Jelly's Gemini QA is built around. JELLY TODAY no — AutopilotScreen.tsx is a status view of internal crons, not user-facing series	Do not build user-facing autopilot; instead offer 'draft next episode' from a series template (see Remix) that lands in the editor for review. components/animate/screens/live/AutopilotScreen.tsx (keep as internal)	L \$0		
Niche Finder with real channel data / trend research TubeGen StoryShort Faceless.so	Research-first buyers want it, but every reviewer says TubeGen's is thin and it is a different product. JELLY TODAY partial — NicheFinderScreen.tsx is a creator-model browser with no live data	Do not build live channel scraping; keep CREATOR_MODELS as opinionated title formulas. components/animate/screens/browse/NicheFinderScreen.tsx, lib/vater/creator-models.ts	L \$\$ data APIs		
Frontier model picker (Veo/Kling/Sora/Seedance per clip) StoryShort Argil Vidnoz Zebracat Faceless.so	Buyers like the menu, but per-clip \$0.50–2 shots blow up the \$1–7 per video promise and break style lock. JELLY TODAY partial — Wan2.2 + Animate2 only, deliberately	Do not add a model menu; keep one locked motion path and let quality be a single 'motion on/off' choice. vater.py post_animate_scene (9265)	M high per clip		

Build order for this month

The ten items in the order they should ship. Nine of ten are **S-effort** and add **\$0** of recurring cost — they surface cost math, engines and routes that already exist. Items 1–3 are the trust triad: show the price, draft cheap, refund failures. Ship those before the first invite goes out. Note: nine of these are NOW rows in the matrix; item 10 (scheduled publish) is tagged NEXT there but rides along because it is a single field on the OAuth path Jelly already uses.

1	☐ Show 'Estimated \$X' before every Generate (stills vs +motion) using check-budget.ts estimateCents in BillingBlock/VisualsStep, plus a per-minute table on PricingScreen.	S \$0 (uses existing cost math)
2	☐ Make stills-draft the default and motion an explicit second paid step (post_animate_all_scenes) so first renders are cheap and QA'd before spend.	S \$0; actually lowers per-video cost
3	☐ Auto-refund failed renders in reconcile-renders.ts and show the reason in the receipt.	S \$0
4	☐ Chapter timestamps + hashtags in DescriptionStep from scene timings.	S \$0
5	☐ Per-scene motion/transition controls (pan/zoom/still, crossfade) in VisualsStep over existing Remotion animations.	S \$0
6	☐ Surface chart/map/header overlay toggles in VisualsStep (engine already in vater.py + remotion/scenes).	S \$0 (Qwen local)
7	☐ 'Remix this project' on ProjectDetail via new-from-style with style, voice, characters, soundtrack pre-filled.	S \$0
8	☐ Pricing calculator on landing/PricingScreen (videos × minutes × motion → dollars) to beat the \$149/mo comparison.	S \$0
9	☐ Bring-your-own narration upload in VoiceoverStep with alignment-driven captions and scene timing.	M \$0 (Whisper/alignment already local)
10	☐ YouTube 'Publish at' scheduling in PublishPanel using publishAt on the existing OAuth path.	S \$0

Ship 1–3 before any invite goes out

Show the estimate, make stills-draft the default, auto-refund failures. Those three turn the receipt from a surprise into a promise — and they are the exact three things Vidnoz (2.3/5), StoryShort (2.5/5) and Zebracat get flamed for missing.

Then 4–10, in any order

Chapters, per-scene motion, overlay toggles, Remix, the pricing calculator, BYO narration and scheduled publish. Each is a single screen. Only BYO narration is M-effort; the rest are an afternoon each over code that already ships.

Zebracat

zebracat.ai

9 ✓ Jelly has

20 ⬢ partial

9 ✗ missing

What they do well

- Published per-feature credit burn table (credits/min per component) with a worked example — makes cost predictable before render; Jelly could show a similar pre-render estimate per step (voice tier, motion on/off, Wan clips).
- Choose visual mode + aspect ratio + duration slider + character-consistency toggle on ONE pre-generation screen, then land straight in the editor — very low-friction 'first video in 2-5 min' path.
- Brand Kit auto-applied to every video (logo, palette, fonts, business name) — one-time setup, zero per-video work.
- Reward Center: 10 credits per external review — converts happy users into G2/Trustpilot/Capterra social proof (1,000+ 5-star reviews claimed).
- Six per-format SEO landing pages/free tools (Shorts, TikTok, Reels, explainer, aspect-ratio calculator) plus a Skool community and public feature-voting board driving organic acquisition.

What users complain about

- Credit bait-and-switch: AppSumo lifetime holders had credits replaced by a 30-videos/month cap plus separate AI-scene credits (AppSumo 3.9/5, 18 one-taco reviews).
- Dual metering (video count + generative-AI credits) is opaque; unused credits expire monthly, no rollover, no partial refunds; refunds described as 'a nightmare'.
- Limited post-generation editing: users report they could not change text/captions/voice/music after render and had to recreate and burn credits; only 3-5 free edits per video.
- Export/render failures ('could not export 50% of the videos'), crashes at peak, 3-10 min render times; support unwilling to refund credits during outages.
- Quality: AI faces distorted, images out of sync with audio, stock clips not matching script, avatars 'slightly robotic', synthetic look on human/lifestyle content.
- Hard 5-minute max and 1080p cap on every tier (no 4K, no true long-form); free tier is 30 s with watermark and no storyboard preview.
- Thin voice library vs. competitors and only basic caption styles (no word-by-word animation); no auto-posting/scheduling and no API below Enterprise (\$599+).

Feature list vs. Jelly

✓ Jelly has it · ⬢ partial · ✗ missing		
✓ Prompt/topic to video	⬢ AI automated editing (cuts, transitions, sync)	
✓ Script to video	⬢ Templates library	
✗ Blog/URL to video	⬢ Output: MP4 up to 1080p, watermark on free	
⬢ Audio/podcast to video	⬢ Direct publishing / scheduling	
✓ AI script generator (standalone + tone)	⬢ Commercial rights + YouTube monetization FAQ	
⬢ Six selectable visual modes at start	✗ API access (Enterprise only)	
✓ AI image scenes in 78 preset styles	⬢ Teams / roles & permissions (Enterprise)	
✓ Custom AI styles	⬢ Credit system with published per-feature burn table	
✓ Character consistency toggle	✗ Reward Center (earn credits for reviews)	
⬢ Generative video scenes via 3rd-party models	✗ Affiliate program (30% recurring)	
⬢ Stock footage + music library (millions of clips)	✗ 7-day money-back guarantee (conditional)	
✗ AI avatars (120+ library, custom, digital twin)	⬢ HelpCat in-app AI chatbot + tiered support	
✓ Voice cloning	⬢ Skool community + public feature voting board	
⬢ 175+ language TTS, standard vs ultra-realistic tiers	✗ Free SEO tool pages (calculators, per-platform generators)	
✗ Video translation with lip-sync (25+ languages)	⬢ Multi-currency + wallets checkout	
✓ Auto captions with styling	⬢ Free tier with no card	
⬢ Aspect ratio 9:16 / 1:1 / 16:9	⬢ Analytics/SEO tools (titles, tags, thumbnails, trend finder)	
⬢ Duration slider		
✗ Brand Kit		
✓ Built-in post-generation editor		
⬢ Metered 'free edits per video' (3/3/5/unlimited)		

SOURCES (26)

www.zebracat.ai www.zebracat.ai/pricing www.zebracat.ai/features www.zebracat.ai/features/text-to-video www.zebracat.ai/features/ai-avatar-generator
www.zebracat.ai/features/ai-scene-generator www.zebracat.ai/features/ai-automated-video-editing help.zebracat.ai
help.zebracat.ai/article/930ed02bf73848769e67211451efe020... help.zebracat.ai/article/e9210101f8ad4466a7d75bdd18865eb...
help.zebracat.ai/article/c75eadbcbaec48bb9cd7a114cb5fc1cf... help.zebracat.ai/article/90a61f36ffd84014a8fb8d82ea2136a8...
help.zebracat.ai/article/6a9b935eb42247c0bd73e5cef2cfdaeb... help.zebracat.ai/article/0de66b0bbe864ddf97bcf877179bc727... +12 more

Faceless.so

Faceless LLC / Side-Products; "faceless channel on autopilot" — shorts generator + 7-platform auto-poster + agent API

4 ✓ Jelly has

16 ● partial

20 ✗ missing

What they do well

- Agent-first surface: llms.txt + llms-full.txt, every doc readable as .md, OpenAPI, hosted MCP with 27 tools, CLI with --json/--wait, and a Claude Skill that installs itself — 'give your agent one scoped key and it runs the channel'. Jelly's planner/QA already runs on LLMs; exposing render/publish/credits as MCP tools is cheap and differentiating.
- Credits charged once at creation, everything downstream free (re-render, captions, publish, schedule, analytics) with model catalog + live costs (GET /options?kind=models) so users/agents pick quality by budget — a clean 3-tier (storyboard / motion_lite / motion_pro) that maps directly onto Jelly's stills-vs-Wan-motion choice.
- 'Review from your phone, polish on desktop': every scheduled episode sits in a queue that's editable until post time; approve on mobile, open timeline editor on desktop.
- No-signup, fully interactive editor demo (/demo) with rendering/upload disabled — lets prospects feel the tool before creating an account.
- Series wizard = 10-minute setup (niche/custom prompt → style → caption theme → voice → 30/60/90s → posting time/days/timezone → destination accounts) then daily autopilot; plus a per-episode 'generate now' button and pause toggle.
- Idempotency-Key on every mutating call (24h replay) so retries can never double-charge credits; 402 insufficient_credits vs 403 usage_limit_reached vs 429 with Retry-After are distinct, documented error types.

What users complain about

- No refunds once a single credit is used ('unable to refund plans once credits have been used'); paid plans also cap total projects (403 usage_limit_reached) on top of credits — two ceilings surprise users.
- 'Sign up free' but no free render tier: their own comparison pages list 'No free tier; entry point is the \$29/month Starter plan' and the public demo disables rendering/upload; pricing shown as \$24 on /pricing but \$29 across compare/templates/integrations pages — inconsistent.
- Editor is admittedly 'workflow-focused, less granular than a general video editor' and 'voice library smaller than ElevenLabs's catalog directly'; scheduling UI called less polished than AutoShorts (calendar-based, no drag-to-reschedule).
- Short-form only (30/60/90s presets, 9:16); no long-form/16:9 pipeline, no thumbnail generator, no stated resolution/4K.
- 'AI slop' / generic-niche risk: reviewers of this category note pure automation plateaus and needs human oversight; the product still leans on Pexels stock B-roll and gameplay backgrounds, so visuals are not unique per channel.
- Trust signals are thin/mixed: very few independent reviews (Groupify 4.75/5 from only 4 reviews; no reachable Trustpilot page for faceless.so; sibling brands faceless.video/facelessreels have 1.8–3.9 TrustScores with refund complaints that spill over in search); ToS governed by Uttar Pradesh, India; ToS forbids 'automated means to access the Service' while selling an API; crypto token/spin-and-win/warmed-account upsells read gray-hat.
- Integrations directory (Zapier/Make/Drive) says pages 'will appear here once they're seeded' — advertised but not delivered; several linked compare URLs 404.

Feature list vs. Jelly

✓ Jelly has it	● partial	✗ missing
✗ Niche/series autopilot (set-and-forget channel)		✗ AI Agent natural-language editing (Growth+)
✗ Research-seeded scripting (Google Trends, Wikipedia, Google News, Reddit, fact grounding)		● Phone review/approve, desktop polish
● Script & hook generation + free script/title/caption tools		✗ Interesting-clips detection per project
✓ Paste your own script (script-first) via UI or API		✗ Brand kits (logos, colors, fonts)
✗ URL imports: Reddit post, blog article, tweet/X thread as source		● Custom templates + unlimited exports, no watermark
● Multiple video formats/creators (Reddit story, fake-text drama, ChatGPT-video, character story/animation, audio-to-visuals, split-screen, streamer-blur, montage, music video)		● Short-form only length presets (30/60/90s), 9:16 vertical output
● Three visual 'models' priced by motion level (storyboard / motion_lite / motion_pro)		● Direct publishing to 7 platforms (YouTube, TikTok, Instagram, Facebook, LinkedIn, Threads, X)
✗ Frontier gen-video creators (Google Veo 3 / 3.1, Seedance v2) as separate tools		✗ Scheduling + content calendar (weeks ahead), per-platform metadata
✗ Background gameplay/footage library (Minecraft parkour, Subway Surfers style) + random background option		✗ Auto-X (Twitter bot) audiogram/faceless posts
✗ Stock B-roll (Pexels) swap in editor		✗ Cross-platform analytics dashboard
✓ Visual styles + masterStyle + globalNegativePrompt		● AI caption generator, hashtags, post ideas & audio trends
✗ AI actors/avatars, face-swap, virtual try-on, UGC videos		● Teams/workspaces with seat limits, invites, folders
● ElevenLabs voices (50+) in 70+ languages		✗ REST API v1 with scoped keys, idempotency keys, rate limits, typed errors, OpenAPI spec
✓ Voice cloning (custom voice library) on every paid plan		✗ MCP server (remote + local, 27 tools), CLI (npm faceless-cli), Claude Skill; llms.txt/llms-full.txt
● Cinematic captions with per-word styling and emphasis detection; caption themes; emoji overlay		✗ Zapier / Make.com / Google Drive integrations
✗ Caption existing footage (video or audio URL) for free		● Register hosted assets by URL
✓ Background music by mood (Suno) + sound effects		● Credits ledger + balance API, top-ups, prorated plan changes, cancel-anytime
● Faceless Studio timeline editor (real-time, effects, zooms, transitions, progress bars, animations, waveforms)		✗ Warmed human-run TikTok/IG accounts (from \$349/mo)
✗ No-signup interactive editor demo		✗ Growth loops: affiliate program, loyalty rewards, spin-and-win credits, \$FACELESS token, AppSumo LTD
		● Support: premium live chat all plans, call support on Ultra; tutorials, 72-post blog, glossary, public roadmap
		● Full ownership + commercial rights; monetization FAQ (YPP, TikTok rewards)

SOURCES (25)

faceless.so faceless.so/faceless faceless.so/pricing faceless.so/llms.txt faceless.so/llms-full.txt faceless.so/developers/docs/credits-and-limits.md faceless.so/developers/docs/videos.md faceless.so/developers/docs/series.md faceless.so/developers/docs/publishing-and-scheduling.md faceless.so/agents faceless.so/warmed-accounts faceless.so/integrations faceless.so/compare faceless.so/compare/autosHORTS-vs-faceless +11 more

Vidnoz AI

vidnoz.com — avatar-first AI video generator + Vidnoz Flex (record/host/track) + Vidnoz Gen (image/video models) + Open API

5 ✓ Jelly has

14 ⚪ partial

16 ✗ missing

What they do well

- Daily-reset free credits (60/day, no card) — every visitor can produce something today; the reset cadence pulls people back daily and converts better than a one-time trial.
- AI Video Wizard intake form: one screen for topic|script|URL|PDF, scene count, aspect ratio, tone, and language — then an editable per-scene script draft with 'recreate' before any render credit is spent.
- Avatar/voice pickers with rich filters (language, gender, age, region, emotion) plus instant preview; reviewers call the filter grid the fastest part of the flow.
- 'AI Image' button on any visual element in the editor — regenerate just that image from a prompt without leaving the scene.
- Hosted share page with brand kit, CTA/polls/forms and per-viewer analytics (watch %, device, geo) — turns each render into a trackable landing page.
- Transparent credit-cost table on pricing page (0.5 credit/sec standard, 2/sec expressive) plus a Credit History ledger page.

What users complain about

- Billing traps: 'billed yearly' fine print, surprise auto-renewals, being charged \$220 when expecting \$19.99, no refunds on monthly plans — Trustpilot ~2.3/5 (~42% one-star) driven almost entirely by billing/cancellation, not product.
- Credits are opaque/expire: no rollover, add-on packs die with the cycle, hard to predict monthly cost; misleading 'FREE' headline when most useful features are paid.
- Lip-sync not visible in preview — you only see it after spending credits on a render, so fixes mean re-rendering and burning minutes.
- Avatar/voice quality: voices sometimes robotic or rushed, lip-sync drifts on fast lines and long scripts, results feel 'AI' (70–80% there).
- Editor is shallow vs pro tools — limited timeline control, few caption styles, slow rendering on free/standard queue at peak times.
- Support is email-only with slow/inconsistent responsiveness; short per-video caps on lower tiers (3 min free, some report 5-min cap on Starter, translator 60–90s free).

Feature list vs. Jelly

✓ Jelly has it	⚪ partial	✗ missing
✓ Script text input (paste script → avatar reads it)	✓ AI Video Wizard: topic/prompt → full video	✓ Auto subtitles / captions with show/hide + style
✗ URL / blog / article → video	✗ PDF / PPT / DOC → video	⚪ Aspect ratios 16:9 / 9:16 / 1:1
✓ AI Script Generator (standalone scripting aid)	⚪ Stock AI voice library (2,000–2,660+ voices, 140+ languages)	⚪ Output limits: 720p+watermark free, 1080p paid, up to 60-min videos
✓ Voice cloning (Business tier or \$9.99/mo add-on)	⚪ Emotional / expressive TTS	✗ Brand Kit
✗ Video translation + AI dubbing with lip-sync	✗ Realistic AI avatars (1,900+ stock, filters, half/circle/square crop)	✗ Team workspace (up to 1,000 seats), comments, SAML/SSO
✗ Talking Photo / Expressive Avatar (photo → talking head w/ emotions & gestures)	✗ Custom AI Avatar / Avatar Pro digital twin	⚪ Hosted share page + embed + video email (Vidnoz Flex)
✗ AI Product Avatar / UGC ad generator	⚪ Templates (2,800–3,400+, landscape+portrait, by industry)	✗ Video analytics (views, watch rate, viewer location/device/source)
⚪ Scene-based canvas editor with layers	⚪ Free stock media library (images, music, icons, text presets)	✗ Screen/webcam recorder + teleprompter (Chrome extension + web)
✓ AI Image inside editor (replace visuals with generated image)	⚪ Vidnoz Gen: image-to-video / text-to-video via 3rd-party models	⚪ Direct social publishing / scheduling
		✗ SCORM export for LMS
		✗ Open API (avatar video, talking photo, TTS, headshots, avatar creation)
		⚪ Credit-metered billing with daily free credits + credit history page
		✗ No-signup try-before-login for some tools
		⚪ My Creations library + Tools Library dashboard
		⚪ Multi-language product UI + how-to help center
		✗ Trust/compliance: ISO/IEC 27001, isolated data center (Enterprise)
		✗ Interactive 24/7 avatars (chat-driven)

SOURCES (19)

www.vidnoz.com www.vidnoz.com/pricing.html www.vidnoz.com/ai-video-generator.html www.vidnoz.com/how-to/index.html
www.vidnoz.com/ai-solutions/clone-my-voice.html www.vidnoz.com/ai-solutions/ai-video-translator-with-lip-...
www.vidnoz.com/get-started/how-does-vidnoz-ai-studio-bill... www.vidnoz.com/refundpolicy.html www.vidnoz.com/docs/authentication.html
www.vidnoz.com/docs/limits.html www.vidnoz.com/video-creation/what-is-vidnoz-flex.html chromewebstore.google.com/detail/vidnoz-flex-video-record...
aiblogfirst.com/vidnoz-ai-pricing www.aitoolspolice.com/reviews/vidnoz +5 more

TubeGen AI

tubegen.ai — Eddie Eizner / TubeGen LTD; faceless-YouTube full-pipeline SaaS (Title → Script → Voiceover → Visuals → Music → Editor → Thumbnail → ...

6 ✓ Jelly has

19 ○ partial

14 ✗ missing

What they do well

- Interactive no-login product demo on the homepage that walks the exact 7 steps with real settings and credit prices (visuals '90 credits', 30 scenes, thumbnail prompt, 3 music segments) — de-risks a no-trial product; Jelly could ship the same with a canned project.
- Price shown at the point of choice: Turbo 20 / Standard 40 / Best 90 credits per image, 'Generate 30 images' total, and a pricing calculator (videos × length × pacing × visual mix) that recommends a plan — mirrors Jelly's per-video receipt but moves it BEFORE render.
- 'Copy channel' / Style presets: paste a YouTube channel URL and the whole pipeline (titles, script tone, art style from pulled frames, voice, language) inherits it in one click; Styles are saved objects snapshotted onto projects.
- Voiceover-as-timeline: narration is generated first and every downstream step (scene split, image timing, per-segment music moods, chapter timestamps in the description) keys off it; editor opens pre-assembled with SSE progress per chunk.
- Thumbnail Studio depth: region edits/expression swaps/background replace stacked in one apply, A/B variants with an AI critic score and side-by-side compare, publish straight to YouTube; plus 'every tool works standalone' + à-la-carte suites from \$27 to lower the entry price.

What users complain about

- High entry price (\$149/mo Starter) with no free trial and no refunds; must buy blind (own FAQ + OutlierKit + tubeaitools reviews).
- Image credits burn: 'Best' images at 90 cr, batch scene generations that fail or come out inconsistent still cost credits; users ask for cheap watermarked previews (Trustpilot 4★/3★, Apr 2026); TubeGen's answer is manual credit refunds via Discord support.
- Consistent characters break down with complex descriptions; visual quality inconsistent across scenes (Trustpilot).
- Scripts read generic and need substantial manual editing; hook/retention optimization is basic (OutlierKit, Grokipedia).
- No keyword research/outlier detection/analytics suite despite 'Niche Finder' branding; reviewers steer research-first buyers elsewhere.
- Interface described as less intuitive than competitors and support responsiveness limited by third-party reviewers; affiliate 1★ (Aug 2026) alleges withheld commissions.
- Monetization/AI-slop risk: YouTube July 2025 'inauthentic content' policy; a 1★ Trustpilot review says AI content can't get monetized; Fortune profile ties the tool to 'AI slop' sleep-history channels.
- Trust signals are thin/lumpy: only 18 Trustpilot reviews, 12 posted on a single day (23 Apr 2026); third-party sites report stale/contradictory prices (\$29–\$99, 2.8/5) creating confusion.
- Editor caps exports at 30 minutes and offers only crossfade + pan/zoom; no shorts/vertical-first workflow, no scheduling, no multi-platform posting.

Feature list vs. Jelly

✓ Jelly has it · ○ partial · ✗ missing	
✗ Interactive no-signup homepage demo	○ Adjustable crossfade transitions
✗ Niche Finder (channel database + real-time trends)	✓ Styled captions
○ AI Title Generator (topic or channel-style)	○ Export up to 30 minutes (MP4), longer via assets folder
○ Retention-structured Script Writer with Copy Style	○ Aspect ratio selection
✗ Multi-language script + voiceover (8 languages)	○ Thumbnail Studio (templates, characters, objects, text, region edits)
○ AI Voiceover library (ElevenLabs + Inworld engines)	✗ Thumbnail A/B variations + AI score + side-by-side compare
✓ Voice cloning (plan-gated)	○ AI Descriptions with hashtags + auto chapter timestamps
✗ Bring-your-own voiceover audio	✓ YouTube OAuth publish (video + thumbnail) and analytics link
✓ Scene Builder: one-pass image generation timed to narration	✗ Team seats
○ Per-image quality tiers with visible credit price	○ Active-project caps + archive
○ Custom Art Styles from reference images or a YouTube channel	✗ Public API v2 + SSE progress streams + docs site
✓ Consistent Characters across scenes and videos, re-rendered in any style	○ Every tool works standalone
✓ Per-scene animation (still → motion clip)	○ Credit system: monthly allotment, 1-cycle rollover, top-ups, usage view
✗ Storyblocks stock B-roll matched to script with mix ratio	✗ Pricing calculator
✗ Auto Overlays (charts, maps, headers)	○ Built-in YouTube Automation Course / Learning Centre
✗ AI Avatars / talking-head presenter (HeyGen-powered)	○ Discord community + support tickets + email
○ AI Background Music scored per segment	✗ Affiliate program
○ Pre-assembled timeline editor	○ Commercial rights + monetization claim
○ Camera effects per scene (pan/zoom, alternating)	✗ Comparison / alternatives / first-party 'honest review' SEO pages
	✗ Scheduling / autopilot / series posting

SOURCES (35)

tubegen.ai tubegen.ai/pricing tubegen.ai/features tubegen.ai/features/ai-script-writer tubegen.ai/features/ai-voiceover
tubegen.ai/features/ai-image-generator tubegen.ai/features/ai-animation tubegen.ai/features/ai-avatars tubegen.ai/features/stock-footage
tubegen.ai/features/auto-overlays tubegen.ai/features/consistent-characters tubegen.ai/features/custom-art-styles
tubegen.ai/features/ai-background-music tubegen.ai/features/video-editor +21 more

NoLang

no-lang.com — Mavericks Inc., Tokyo (Univ. of Tokyo alumni); launched Jul 2024; ~200K registered users, 100–500+ corporate clients; ISO 27001/9001;...

2 ✓ Jelly has

21 ○ partial

14 ✗ missing

What they do well

- Watch page keeps the just-generated video as a pinned 'reference video' — you type a delta ('slower narration', 'more pop', 'same structure, new topic') and get a sibling video that inherits design/structure/voice; any past video can seed a series. Cheap way to give Jelly 'remix from this project'.
- One-box prompt-first home with only two side choices (aspect ratio incl. 'auto', duration incl. 'auto'); everything else lives behind 'Saved settings' presets that power users (and the API) reuse — nice split between fast lane and power lane.
- Failed-generation credits auto-refund within ~20 min, credits roll over, and the API pricing page publishes an exact credit table per minute (image-gen ON/OFF) plus auto-recharge with a monthly hard limit — copy the transparent per-minute table and the hard-limit safety UX.
- Real-avatar flow renders a static-avatar draft first, lets you finish all edits, then a separate 'Animate' button spends the expensive credits (with a warning if the video isn't updated) — a good pattern for gating Jelly's costly Wan2.2 motion pass.
- Editor quality-of-life: 'apply interval to all scenes', in-script pause markers, cross-video pronunciation search + dictionary, and AI-suggested BGM/SFX/background per script instead of keyword search.
- Curated sample gallery by content category (trivia, news, yukkuri history) with playable examples that become templates, plus per-genre tutorials built entirely with the product itself.

What users complain about

- Free tier is nearly a demo: 3 generations/month, watermark, videos disappear after 2 weeks; most wanted features (avatars, no watermark, real avatar, image gen) are paywalled.
- Generic, formulaic visuals — AI-picked stock/CC images feel unoriginal; videos look 'one-dimensional' vs human-made; fixed transitions, hard to place logos or fine-tune timing/audio fades.
- TTS mispronunciations and mechanical intonation in Japanese (e.g., 山積み read wrong); non-Japanese output quality noticeably weaker; scripts get unintentionally summarized or misinterpreted on long inputs.
- PC-only editing/download, no mobile app; Chrome-extension videos and CC-BY-SA attribution requirement limit commercial use; default-public sharing setting is a privacy trap.
- Copyright: royalty-free BGM triggers YouTube false-positive claims; stock images can later prove unlicensed; no refunds for unused subscription/credits and credits vanish on cancel.
- Discoverability: no Trustpilot/G2/Product Hunt footprint (0 English reviews found); pricing only in JPY; auto-post lacks scheduling; API is a separate prepaid wallet from app credits and API videos expire after 2 weeks.

Feature list vs. Jelly

✓ Jelly has it · ○ partial · ✗ missing	
● Prompt-first generation (Aug 12 2026 'new generation')	✓ Image generation AI toggle (quality tier)
● Iterative 'reference video' remix loop	● Stock/CC visuals + royalty-free BGM library + SFX + background videos
● Official sample gallery as templates	● AI recommendations for BGM/SFX/background video
● Script mode (verbatim read, single or multi-speaker)	● New editor (May 2026): character + preview/lane model
● Dialogue ('yukkuri'-style) two-character videos	✗ Tempo controls: apply interval to all scenes + in-script pauses
✗ PDF/PPTX presentation-to-video	✗ Text style presets (gradients, shadows, neon), shapes and masks
● Audio/video-to-video (transcribe + captions + avatar)	● Aspect ratios 16:9 / 9:16 / 1:1 + post-hoc conversion
● Image-set to video	● Length caps and output
✗ Chrome extension: web page → video	● SNS auto-posting (YouTube, Instagram, TikTok, Threads, Facebook)
● Saved settings (reusable generation presets)	● Share links, public gallery, privacy default
● Voice library 300+ voices with search/filter	✗ Slide generation (3 modes) + slide → video handoff
✓ Voice clone from ~1-min in-app recording	✗ Public REST API + OEM/white-label
✗ Pronunciation correction + dictionary	● Credit UX: refresh on signup anniversary, rollover, auto-refund on failure
✗ Multilingual output (18–20+ languages) + one-click language conversion	✗ Attribution / CC-BY-SA licensing requirement
✗ Multilingual subtitle dubbing of uploaded media	✗ Team/asset sharing + dedicated support (Business)
✗ Avatar library 100+ (Live2D, PSD standing art, 3D 'real')	● Referral/affiliate program
● Photo → character avatar ('写真をキャラ化')	● Support channels + status transparency
✗ Real avatar / 2D animation of uploaded image (Premium/Business)	● Script export + prompt saving + video management
✗ Custom Live2D / PSD avatar upload + Live2D motion mapping	

SOURCES (42)

no-lang.com no-lang.com/en no-lang.com/pricing no-lang.com/en/pricing no-lang.com/faq no-lang.com/business no-lang.com/api-docs
no-lang.com/api-docs/pricing no-lang.com/api-docs/reference no-lang.com/tutorial no-lang.com/tutorial/updates/2026-08-12-new-generation
no-lang.com/tutorial/updates/2026-05-31-new-editor no-lang.com/tutorial/basic/generate-video no-lang.com/tutorial/basic/generation-methods +28 more

Argil

argil.ai — Paris/YC-backed AI clone-avatar short-form video platform, mid-pivot (Aug 2026) to an agentic "Director / Playground / Templates"...

2 ✓ Jelly has

25 🟡 partial

12 ✗ missing

What they do well

- Preview-before-you-pay: full script/voice/B-roll/caption preview with a still avatar; credits are only burned when you hit Generate (lipsync). Jelly could show a free 'animatic' (stills + TTS) before charging for motion.
- Per-clip B-roll 'shuffle' (3 free re-rolls) with Display front/behind and Length 1/2 toggles — a zero-prompt way to fix a bad shot; plus 'apply to all medias' for layout changes.
- Transparent credit table (video/voice/B-roll/royalty per unit) with a worked 30-second example, PAYG credits that never expire, and Auto Top-Up with a monthly spend cap — reduces billing anxiety while keeping usage-based feel.
- 'Remix' on any generated clip (reuse prompt + voice + first frame) and everything auto-saved to Assets so generations become reusable B-roll — compounding library value.
- Editing-tips doc baked into the flow: 'split sentences to add a zoom', 'shorter clips', 'record your voice instead of typing' — opinionated guidance that raises output quality without more features.
- Docs published as Mintlify + llms.txt/OpenAPI, and model-side rejections (Seedance/Veo moderation) auto-refund credits with a plain-language reason — trust signals worth copying for a receipts-based product.

What users complain about

- Lip-sync is 'a bit off' / uncanny at times; reviewers advise hiding it with B-roll every 5–7 s; blinks and repetitive hand gestures clip through the body; static camera.
- Trustpilot 2.8/5 (3 reviews): advertised 5-day trial hits a paywall immediately after signup; annual subscriber refused any pro-rata refund on cancellation; no G2 listing.
- Subscription required (no true pay-per-video); credit math is confusing (marketing '25 min for \$39' vs docs 160 credits/min Atom ⇒ ~9 min), plan names/limits differ across pages (Classic/Pro/Scale vs new Pro/Business/Enterprise; docs still show old plans).
- Watermark on free/trial output, premium models (Seedance 2.0) and clips locked during trial, billing geo-restricted ("Payments not available in your region").
- Quality complaints on stock avatars: 'low-effort, unconvincing actors', output 'leans corporate'; not suitable for 10+ minute long-form; each outfit change = new avatar/style; render 5–15 min (Pro faster).
- Strict content moderation on licensed avatars (finance/crypto/politics banned) that is 'too strict and sometimes doesn't work'; product-direction whiplash (site now pitches politics/cinema, legacy docs still clone-avatar shorts).

Feature list vs. Jelly

✓ Jelly has it	🟡 partial	✗ missing
✗ AI clone avatar from your own footage (Personal avatar)	🟡 AI Influencer / synthetic avatar builder	✗ Batch/cadence workflow (queue many scripts, hook/A-B variants)
✗ 100+ licensed stock avatars (public + Pro), horizontal & vertical, tag filters	🟡 Avatar Styles (same face, new outfit/scene) + Vary / Use settings / Copy the style of an image	✗ Make (Integromat) automation tutorial + Argil CLI + connectors
✓ Script-first creation (write/paste script or prompt) with paragraph = clip model	🟡 Voice-to-video: upload audio or record in-app, then voice transformation	🟡 Fictions → 'Avatar Action' (Veo 3 / Veo 3.1 Fast / Hailuo / Nano Banana 8-second clips with your avatar)
🟡 Article-to-video (URL → script + pulled media)	✗ Prompt-a-voice (text-described voice design)	🟡 Story — multi-scene mini-film/ad from one prompt (Kling 3.0 / 3.0 Pro)
✓ Voice cloning (20 s–5 min upload) + generated 'voices for your face'	🟡 Bring-your-own ElevenLabs / MiniMax voices (Pro voices) + per-generation model & settings override	✗ Motion Control (transfer body movement from a reference video to avatar photo)
✗ Prompt-a-voice (text-described voice design)	🟡 ~30+ languages via ElevenLabs; captions in any language	✗ Product Interaction / Product Presentation / Product Visual / Animate an Image
🟡 Magic editing pre-sets at creation (voice, captions, B-roll type, layout, music) → pre-edited video in one click	🟡 Auto B-roll from 4 sources: AI image, Getty stock video, Google Images, upload (+ Avatar Action clips)	🟡 Playground — model-by-model image/video/audio generation (Seedance 2.0/2.5/1.5, MiniMax H3 2K, Sora 2 Pro, Veo 3.1, Kling, Lyria 3)
✗ Layouts / avatar positioning (full, PIP corner, split top/bottom or left/right, background)	🟡 Camera angles / zooms / body-language (gesture) per clip	🟡 Director — agentic chat 'idea to finished video' with Agent Plan, storyboard, shots, Agent Skills
🟡 Caption styling (style presets, size S/M/L, position top/middle/bottom) + word-level typo fix + subtitle export	🟡 Background music library with preview & volume	🟡 Templates gallery (Playground + Director recipes)
🟡 Studio editor with preview-before-render (lipsync generated only at final render)	🟡 Aspect ratios 9:16 and 16:9 (1:1 claimed in blog), multi-aspect render from one script	🟡 Assets library + folders + video-project folders
🟡 Length / watermark limits	🟡 Direct publishing to TikTok (OAuth) with Branded-Content disclosure	✗ Team workspaces (seats, invitations, member management, audit logs on Enterprise)
		✗ Multi-language UI (Display Language setting) + iOS app
		🟡 Billing UX: credits balance/renewal date, PAYG credits that never expire, Auto Top-Up with threshold + monthly cap, usage CSV, cost-per-feature table in-app
		🟡 Free tier / trial: 3 free generations, daily credits with monthly cap, watermark, 5-day trial that converts
		🟡 Content moderation policy tiers by avatar type + auto-refund + email appeal
		🟡 Support & community: email, Typeform, Discord, YouTube/Loom tutorials embedded in docs, affiliate program
		✗ Solutions verticals: Creators, Politics & campaigns, Media, Cinema

SOURCES (43)

www.argil.ai www.argil.ai/product www.argil.ai/pricing www.argil.ai/templates www.argil.ai/solutions www.argil.ai/blog www.argil.ai/blog/runway-pricing-plans
www.argil.ai/blog/how-do-leading-ai-avatar-services-compa... www.argil.ai/blog/how-to-upload-a-youtube-short docs.argil.ai/llms.txt docs.argil.ai/resources/subscription-and-plans
docs.argil.ai/resources/pay-as-you-go-pricings docs.argil.ai/resources/api-pricings docs.argil.ai/resources/create-a-video +29 more

StoryShort AI

storyshort.ai

9 ✓ Jelly has

20 ⚪ partial

9 ✗ missing

What they do well

- Credit-cost table on the pricing page ("What do credits actually make?") plus 'Estimated cost: N credits' before every generate and a per-video metadata dialog showing script + credits used — cost is legible at every step.
- Series = 'set topic, style, cadence once' with pause/resume, calendar view, and every queued video openable in the full editor before it auto-posts (edit-then-let-schedule-resume).
- Inspiration tab: breakout videos in your niche with a one-click 'analyze then make an original with the same hook style' prefill into the create wizard.
- Media picker inside the editor pulls Pexels/Giphy/Google Images/website-scroll capture/AI gen in one panel; audio panel lets you paste a YouTube URL to extract a track and upload MP3/WAV.
- API docs page with a 'Copy for LLM' button and webhook-first flow; Zapier/Make connect via a generic 'Custom webhooks' publishing target.
- Model picker with plain-language tiers and price per clip ("Fast, affordable clips (5 credits/clip)" vs 'Best cinematic quality, characters can speak (15 credits/clip)') and automatic credit refunds on moderation/failed renders.

What users complain about

- Support silence: Trustpilot reviewers report no reply to emails, no complaint path, no refund; affiliate says >\$1,000 commissions marked paid then reversed (Trustpilot 2.5/5).
- Output quality below the marketing examples and generation times 'far longer than promised' (Trustpilot); reviewers cite visual inconsistency between scenes and 'visibly artificial' AI images.
- No refunds, no PayPal, subscription-only; free account can only preview (export watermarked/locked), 10 free credits — 'restrictive free plan' and 'pricier than competitors'.
- Limited manual control: basic transitions/effects, prompt-to-video accepts image but not video references, editor learning curve; long-form and UGC gated to \$129+/mo plans.
- Music library copyright risk (some tracks flagged as likely copyrighted) requiring manual checks; occasional struggles with specialized terminology needing human review.
- Marketing inconsistencies erode trust: FAQ says free tier has 'no watermark' while app gates 'Export watermark-free HD'; refund page still says 'contact us at [your contact email]'.

Feature list vs. Jelly

✓ Jelly has it · ⚪ partial · ✗ missing		
✓ Prompt-to-video (topic/idea input)	✗ Media browser: Pexels stock, Giphy stickers, Google Images, uploads, website capture	
✓ Script-to-video (paste script)	⚪ Caption styling (fonts, layouts, emojis, colors, position)	
⚪ URL / article / Reddit / PDF / video inputs	⚪ Music: library of 1000s, sound effects, upload own, extract from YouTube URL, AI music generation	
✗ AI Agent (chat-based creation)	⚪ Aspect ratios 9:16 / 1:1 / 16:9 (switch in editor)	
✓ AI script writer with random idea + long-form script generator	✓ Thumbnail generation/upload	
⚪ Voice library (100+ voices, ElevenLabs + OpenAI)	⚪ Watermark toggle / HD export	
✓ Voice cloning	⚪ Direct publishing to TikTok, YouTube (Shorts), Instagram Reels	
⚪ Multi-language generation (30-50+ languages) + 5-language UI	✗ Series (autopilot channels) + Calendar scheduling	
⚪ Media type choice: AI images / AI video clips / stock footage / Google Images	✗ Bulk generation (number of videos)	
✓ Image style presets (12+) + Basic/Premium/Ultra image models	✗ Stats dashboard (cross-platform analytics)	
⚪ Frontier video-model picker (Veo 3.1, Kling v3 Pro, Seedance 2.0, Veo 3, Wan 2.2, P-Video 720p/1080p, Grok Imagine, Sora 2)	✗ Inspiration: Discover / Watchlist / Ideas (trend finder)	
✓ Animate first segments / Animate Image (image-to-video motion)	✗ Public REST API + webhooks + Zapier/Make	
✓ Consistent characters + Character library	⚪ Credit transparency + video metadata	
✗ Talking Avatar / UGC avatar / Lip Sync videos	⚪ Library with favorites, share link, publish, download	
✗ Motion Control (video-to-image motion transfer)	⚪ Standalone AI Image / AI Video generators	
⚪ Workflow templates (21+ formats)	⚪ In-app changelog ('What's New' modal) + onboarding checklist/survey	
⚪ Long-form videos (10-30 min) 16:9	⚪ Support: Crisp live chat, in-app help, contact page	
⚪ Full timeline editor (drag/trim tracks)	⚪ Billing UX: Stripe portal, retention offers, promo codes	
✓ Per-segment voiceover text edit + audio regenerate	⚪ Trust / policies	

SOURCES (35)

storyshort.ai storyshort.ai/pricing storyshort.ai/_next/static/chunks/app/(marketing)/pricing... storyshort.ai/ai-long-form-video-generator storyshort.ai/motion-control storyshort.ai/automate-ai-ugc storyshort.ai/api-docs storyshort.ai/tools storyshort.ai/alternatives storyshort.ai/community storyshort.ai/affiliate storyshort.ai/refund storyshort.ai/terms storyshort.ai/sitemaps/workflows.xml +21 more